

Outline

This presentation covers the following topics

➤ Cilk and the Lucata API

- *Basic programming*
- *Data distribution*

*Slides originally developed by
Janice McMahon, Lucata
Corporation*





Basic Programming

Parallelism, spawning/syncing threads with Cilk

Spawning Threads with Cilk

- Specifies function may run in parallel with caller
 - Child thread spawned to execute function and parent continues in parallel w/child
 - Otherwise, parent executes a standard function call
- Spawn location determines location of
 - Synchronization structure
 - Stack frame (if needed)
- Spawn destination
 - Special functions denote spawn location
 - If no direction is given, then spawn is local

```
long f = cilk_spawn fib(a);
```



Thread Synchronization with Cilk

- Current function cannot continue past the `cilk_sync` until all children have completed
- Last thread to reach the `cilk_sync` continues execution
 - no waiting
- Implicit sync at termination of a function

void `cilk_sync`;



Example: Fibonacci

```
#include <stdio.h>
#include <stdlib.h>
#include <cilk/cilk.h>
```

Cilk include file

```
long result;
long fib(long n)
{
    if (n < 2) return n;
    long a = cilk_spawn fib(n-1);
    long b = cilk_spawn fib(n-2);
    cilk_sync;
    return a + b;
}
```

Spawn a thread
for each of the
fib() calls

Wait for threads to
complete to ensure
a and b are valid

```
int main(int argc, char **argv)
{
    long input = atol(argv[1]);
    result = fib(input);
}
```

➤ Recursive function with dynamic spawning

➤ Number of threads depends on data size

- Second spawn could be function call



Sample Program Execution: fib.c

```
>>>>>>> /usr/local/emu/bin/emu-cc -o fib.mwx fib.c
>>>>>>> /usr/local/emu/bin/emusim.x -- fib.mwx 10

SystemC 2.3.3-Accellera --- Mar 24 2021 16:05:40
Copyright (c) 1996-2018 by all Contributors,
ALL RIGHTS RESERVED
Start untimed simulation with local date and time= Thu Mar 25 13:02:34 2021

End untimed simulation with local date and time= Thu Mar 25 13:02:34 2021

>>>>>>> more fib.cdc
*****
Program Name/Arguments:
fib.mwx
10
*****
Simulator Version: 21.3.24
*****
Configuration Details:
Ring Model = Stratix: 3 GC Clusters, 4 MSPs, 1 SRIO
Number of Nodes=1
Total Memory (in GiB)=64
Logical MSPs per Node=1
Log2 Memory Size per MSP=36
GC Clusters per Node=3
GCs per Cluster=8
Uplinks used in rings=false
*****
*****
Simulator wall clock time (seconds): 0
>>>>>>> more fib.vsf
Node ID: Outbound Migrations, Threads Created, Threads Died, Spawn Fails
0: 0, 88, 89, 0

NodeID.MspID: num_reads, num_writes, num_rmws
0.0: 551, 1858, 320

>>>>>>>
```

Compile and run in
simulator; default
configuration is one node

Configuration file with
execution summary

Verbose file with
thread and memory
operation counts

No
migrations;
all data on
Node 0

Extra Credit Question:
Why is #died one more
than #spawned?
(Answer at end of section)



Example: SAXPY with spawn loop

```
#include <stdio.h>
#include <stdlib.h>
#include <cilk/cilk.h>
void saxpy(long n, float a, float *x, float *y)
{
    for (long i = 0; i < n; i++)
        y[i] += a * x[i];
}
int main(int argc, char **argv)
{
    long nth = atol(argv[1]); // number threads
    long size = atol(argv[2]); // array size
    float aval = atof(argv[3]); // constant
    float *x = malloc(size * sizeof(*x));
    float *y = malloc(size * sizeof(*y));
    for (long i = 0; i < size; i++) {
        x[i] = i; y[i] = 0;
    }
    long grain = size / nth; // elts per thread
    for (long i = 0, j = 0; i < nth; i++, j += grain)
        cilk_spawn saxpy(grain, aval, &x[j], &y[j]);
    cilk_sync;
}
```

Allocated on
stack at
current node

Grain size
determined
by number
of threads

Spawn a thread for each section of array

- Iterative loop with static spawning
- Number of threads fixed in code
- Work per thread depends on data size



Sample Program Execution: saxpy.c

```
>>>>>>> /usr/local/emu/bin/emu-cc -o saxpy.mwx saxpy.c
>>>>>>> /usr/local/emu/bin/emusim.x -- saxpy.mwx 4 32 5.0

SystemC 2.3.3-Accellera --- Mar 24 2021 16:05:40
Copyright (c) 1996-2018 by all Contributors,
ALL RIGHTS RESERVED
Start untimed simulation with local date and time= Thu Mar 25 14:02:29 2021

End untimed simulation with local date and time= Thu Mar 25 14:02:29 2021

>>>>>>> more saxpy.cdc
*****
Program Name/Arguments:
saxpy.mwx
4
32
5.0
*****
Simulator Version: 21.3.24
*****
Configuration Details:
Ring Model = Stratix: 3 GC Clusters, 4 MSPs, 1 SRIO
Number of Nodes=1
Total Memory (in GiB)=64
Logical MSPs per Node=1
Log2 Memory Size per MSP=36
GC Clusters per Node=3
GCs per Cluster=8
Uplinks used in rings=false
*****
Simulator wall clock time (seconds): 0
>>>>>>> more saxpy.vsf
Node ID: Outbound Migrations, Threads Created, Threads Died, Spawn Fails
0: 0, 4, 5, 0

NodeID.MspID: num_reads, num_writes, num_rmws
0.0: 72, 47, 21
>>>>>>>
```

Number of
threads
fixed in
code



Parallel Loops in Cilk

- Divides loop among parallel threads, each containing one or more contiguous loop iterations
- Max number of iterations in each chunk is grainsize
- Best for situations where:
 - Threads are spawned locally
 - Work per element is fairly uniform

```
#pragma cilk grainsize = 4  
cilk_for;(long i=0; i<SIZE; i++)  
    {...}
```



Example: SAXPY with parallel loop

```
#include <stdio.h>
#include <stdlib.h>
#include <cilk/cilk.h>
int main(int argc, char **argv)
{
    long size = atol(argv[1]); // array size
    float aval = atof(argv[2]); // constant
    float *x = malloc(T * sizeof(*x));
    float *y = malloc(T * sizeof(*y));
    for (long i = 0; i < T; i++) {
        x[i] = i; y[i] = 0;
    }
    #pragma cilk grainsize = 8
    cilk_for (long i = 0; i < size; i++) {
        y[i] += aval * x[i];
    }
}
```

Grain size determined
by pragma

Use parallel
construct for
spawn loop

- Single thread per loop iteration unless grainsize is specified
- Threads are implicitly synchronized at end of loop



Sample Program Execution: saxpy_for.c

```
>>>>>>> /usr/local/emu/bin/emu-cc saxpy_for.c -o saxpy_for.mwx
>>>>>>> /usr/local/emu/bin/emusim.x -- saxpy_for.mwx 32 5.0

SystemC 2.3.3-Accellera --- Mar 24 2021 16:05:40
Copyright (c) 1996-2018 by all Contributors,
ALL RIGHTS RESERVED
Start untimed simulation with local date and time= Thu Mar 25 14:56:29 2021

End untimed simulation with local date and time= Thu Mar 25 14:56:29 2021

>>>>>>> more saxpy_for.cdc
*****
Program Name/Arguments:
saxpy_for.mwx
32
5.0
*****
Simulator Version: 21.3.24
*****
Configuration Details:
Ring Model = Stratix: 3 GC Clusters, 4 MSPs, 1 SRIO
Number of Nodes=1
Total Memory (in GiB)=64
Logical MSPs per Node=1
Log2 Memory Size per MSP=36
GC Clusters per Node=3
GCs per Cluster=8
Uplinks used in rings=false
*****
Simulator wall clock time (seconds): 0
>>>>>>> more saxpy_for.vsf
Node ID: Outbound Migrations, Threads Created, Threads Died, Spawn Fails
0: 0, 3, 4, 0

NodeID.MspID: num_reads, num_writes, num_rmws
0.0: 99, 73, 44
>>>>>>>
```

**Fewer
memory
operations;
more efficient**

Extra Credit Question:

Why is there one fewer thread than spawn loop?
(Answer at end of section)



Unit Summary: Basic Parallelism

- Dynamic vs. static spawning
- Using Cilk_spawn and Cilk_sync for multi-threading
- Using Cilk_for with grain size to parallelize loop iterations
- Using compiler to build executables
- Using simulator to gather basic statistics and thread counts
- Using basic simulator statistics to verify that correct numbers of threads are spawned

Exercises:

Try fib with different values, measure and verify thread counts
Insert printf to verify results
Try different grain sizes, thread counts

Extra Credit Answers:

Total thread count includes initial thread spawned from the SC
Parallel loop re-uses initial thread for computation





Data Distribution: Part I

Distributed allocation using memoryweb library, targeted thread spawn

Cyclic Array Distribution

- Array of N long integers striped across nodes round robin
- Accessed using 1D array notation (e.g. A[i])
- ONLY works for 64-bit types (e.g. long)

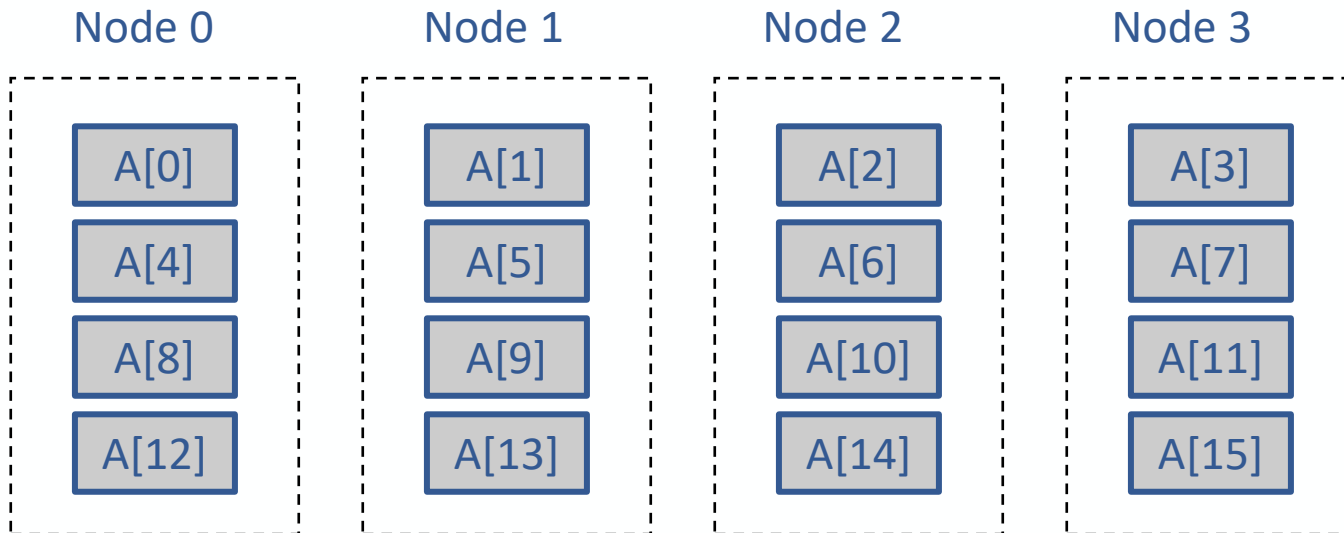
```
long *A = mw_malloc1dlong(unsigned long N);
```



Example: 1D Array

```
#define N 16  
long * A = (long *) mw_malloc1dlong(N);  
for (long i=0; i<N; i++) A[i] = i;
```

Array of N elements, striped
across nodes round-robin



Example: Distributed Vector $AX+Y$

```
#include <stdlib.h>
#include <cilk/cilk.h>
#include <memoryweb.h>
void saxpy(long n, long a, long *x, long *y)
{
    for (long i = 0; i < n; i++)
        y[i] += a * x[i];
}
int main(int argc, char **argv)
{
    long nth = atol(argv[1]); // number threads
    long size = atol(argv[2]); // array size
    long aval = atol(argv[3]); // constant
    long *x = mw_malloc1dlong(size);
    long *y = mw_malloc1dlong(size);
    for (long i = 0; i < size; i++) {
        x[i] = i; y[i] = 0;
    }
    long grain = size / nth; // elts per thread
    for (long i = 0, j = 0; i < nth; i++, j += grain)
        cilk_spawn saxpy(grain, aval, &x[j], &y[j]);
    cilk_sync;
}
```

Memoryweb
include file

Cyclic
distribution of
input arrays

- Code is the same except for initial data allocation
- Array access in program automatically finds correct element in correct node



Sample Program Execution: saxpy_1d.c

```
>>>>>>> /usr/local/emu/bin/emu-cc saxpy_1d.c -o saxpy_1d.mwx
>>>>>>> /usr/local/emu/bin/emusim.x --total_nodes 4 - saxpy_1d.mwx 4 32 5

SystemC 2.3.3-Accellera --- Mar 24 2021 16:05:40
Copyright (c) 1996-2018 by all Contributors,
ALL RIGHTS RESERVED
Start untimed simulation with local date and time= Thu Mar 25 22:40:38 2021

End untimed simulation with local date and time= Thu Mar 25 22:40:38 2021

>>>>>>> more saxpy_1d.vsf
Node ID: Outbound Migrations, Threads Created, Threads Died, Spawn Fails
0: 8, 4, 5, 0
1: 8, 0, 0, 0
2: 8, 0, 0, 0
3: 8, 0, 0, 0

NodeID.MspID: num_reads, num_writes, num_rmws
0.0: 28, 25, 11
1.0: 70, 67, 27
2.0: 44, 26, 74
3.0: 9, 31, 1
```

**Compile and
run in
simulator
with 4 nodes**

**All threads spawned
in node 0, migrate
to node with data**

- Data distributed across nodes causes thread migrations
- All threads spawned in node 0
- Access is parallel, but thread spawning is sequential



Lucata Cilk Extensions for Thread Placement

- Migrate hint specifies node for next `cilk_spawn` operation
 - Argument `p` is a pointer into destination node's memory

`cilk_migrate_hint(p);`

- Directed spawn function combines migrate hint and Cilk spawn into a macro for single-command spawn
 - Implemented as C macro; may require braces for correct operation

`cilk_spawn_at(p) fib(a);`




Example: Distributed Vector and Spawn

```
#include <stdlib.h>
#include <cilk/cilk.h>
#include <memoryweb.h>
void saxpy(long n, long a, long *x, long *y)
{
    for (long i = 0; i < n; i++)
        y[i] += a * x[i];
}
int main(int argc, char **argv)
{
    long nth = atol(argv[1]); // number threads
    long size = atol(argv[2]); // array size
    long aval = atol(argv[3]); // constant
    long *x = mw_malloc1dlong(size);
    long *y = mw_malloc1dlong(size);
    for (long i = 0; i < size; i++) {
        x[i] = i; y[i] = 0;
    }
    long grain = size / nth; // elts per thread
    for (long i = 0, j = 0; i < nth; i++, j += grain) {
        cilk_migrate_hint(&y[j]);
        cilk_spawn saxpy(grain, aval, &x[j], &y[j]);
    } cilk_sync;
}
```

Directed
spawn

- Code is the same except for additional migrate hint
- Equivalent alternative replaces two function calls:

```
cilk_spawn_at (&y[j])
    saxpy(grain, aval,
          &x[j], &y[j]);
```



Sample Program Execution: saxpy_1d_hint.c

```
>>>>>>> /usr/local/emu/bin/emu-cc saxpy_1d_hint.c -o saxpy_1d_hint.mwx
>>>>>>> /usr/local/emu/bin/emusim.x --total_nodes 4 - saxpy_1d_hint.mwx 4 32 5

SystemC 2.3.3-Accellera --- Mar 24 2021 16:05:40
Copyright (c) 1996-2018 by all Contributors,
ALL RIGHTS RESERVED
Start untimed simulation with local date and time= Thu Mar 25 22:40:38 2021

End untimed simulation with local date and time= Thu Mar 25 22:40:38 2021

>>>>>>> more saxpy_1d_hint.vsf
Node ID: Outbound Migrations, Threads Created, Threads Died, Spawn Fails
0: 8, 4, 5, 0
1: 8, 0, 0, 0
2: 8, 0, 0, 0
3: 8, 0, 0, 0

NodeID.MspID: num_reads, num_writes, num_rmws
0.0: 32, 25, 11
1.0: 70, 67, 27
2.0: 36, 26, 82
3.0: 9, 31, 1
```

All threads spawned
in node 0, migrate
to node with data

Extra Credit Question:
Why are all threads still
spawned in Node 0?
(Answer at end of section)



Block Array Distribution

- Array of N pointers striped across nodes round robin
- Each points to co-located memory block of size S
- May be an array of pointers to any type

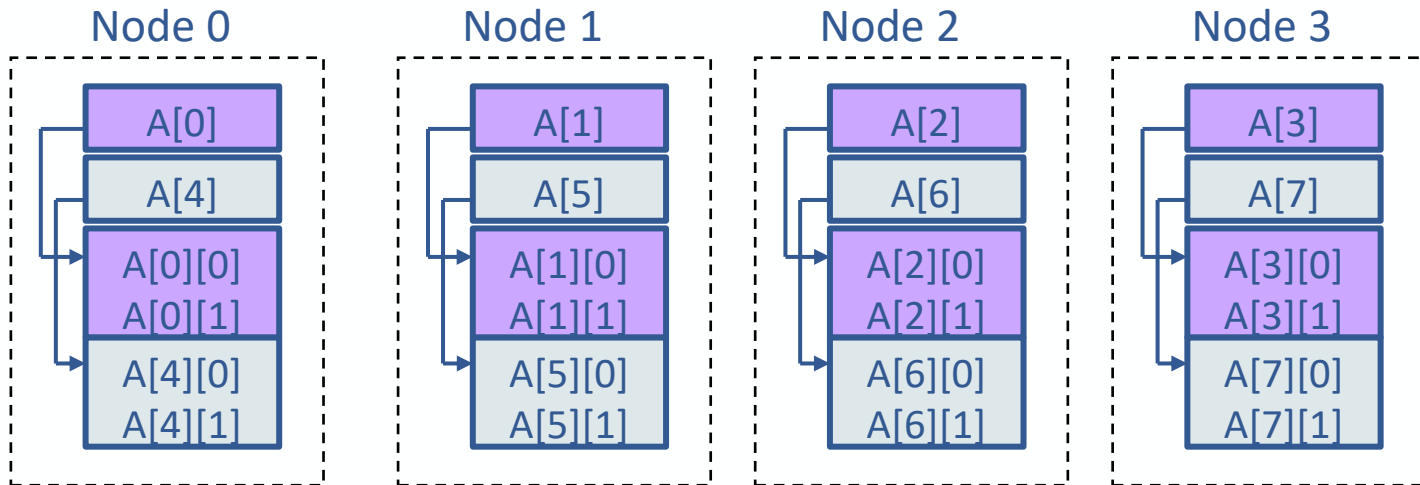
```
void ** mw_malloc2d(unsigned long N,  
                    size_t S);
```



Example: Wrapped Block Distribution

```
#define N 16
long B = 8;
long E = N/B;
long ** A = (long **)
    mw_malloc2d(B, E * sizeof(long));
for (long i=0; i<N; i++)
    A[i/E][i%E] = i;
```


Array of B blocks each with
E elements, striped across
nodes round-robin



Example: Distributed Array of Structures

```
#define N 8
struct element { long a; long b; };
struct element ** E = (struct element **)
    mw_malloc2d(N, sizeof(struct element));
for (long i=0; i<N; i++) {
    E[i]->a = i;
    E[i]->b = 0;
}
```

Array of N blocks each with
containing struct E, striped
across nodes round-robin



Accessing Distributed Data

➤Inputs:

- Array allocated with `mw_malloc2d`
- Index for first dimension
- Number of blocks used in `mw_malloc2d`
- Blocksize used in `mw_malloc2d`

➤Returns address of `A[i][0]`

```
long * =  
mw_arrayindex(void *A, unsigned long i,  
              unsigned long N, size_t S);
```



Example: SAXPY with block distribution

```
#include <stdlib.h>
#include <cilk/cilk.h>
#include <memoryweb.h>
void saxpy(long n, float a, float *x, float *y)
{
    for (long i = 0; i < n; i++)
        y[i] += a * x[i];
}
int main(int argc, char **argv)
{
    long num = atol(argv[1]); // number blocks
    long size = atol(argv[2]); // block size
    float aval = atof(argv[3]); // constant
    float **x = mw_malloc2d(num, size * sizeof(*x));
    float **y = mw_malloc2d(num, size * sizeof(*y));
    for (long j = 0; j < num; j++)
        for (long i = 0; i < size; i++) {
            x[j][i] = j * size + i; y[j][i] = 0;
        }
    for (long i = 0; i < num; i++) {
        cilk_spawn_at (y[i]) saxpy(size, aval, x[i], y[i]);
    }
    cilk_sync;
}
```

- 2D Block allocation
- Number of threads is number of blocks
- Work per thread is block size

Spawn a thread
for each block

Braces required
for macro in loop



Sample Program Execution: saxpy_2d_at.c

```
>>>>>>> /usr/local/emu/bin/emu-cc saxpy_2d_at.c -o saxpy_2d_at.mwx
>>>>>>> /usr/local/emu/bin/emusim.x --total_nodes 4 - saxpy_2d_at.mwx 4 8 5.0

SystemC 2.3.3-Accellera --- Mar 24 2021 16:05:40
Copyright (c) 1996-2018 by all Contributors,
ALL RIGHTS RESERVED
Start untimed simulation with local date and time= Thu Mar 25 22:40:38 2021

End untimed simulation with local date and time= Thu Mar 25 22:40:38 2021

>>>>>>> more saxpy_2d_at.vsf
Node ID: Outbound Migrations, Threads Created, Threads Died, Spawn Fails
0: 12, 1, 5, 0
1: 5, 1, 0, 0
2: 5, 1, 0, 0
3: 5, 1, 0, 0

NodeID.MspID: num_reads, num_writes, num_rmws
0.0: 130, 124, 22
1.0: 90, 104, 31
2.0: 73, 53, 30
3.0: 27, 93, 4
```

Threads spawned
on correct node




Allocate Co-located Data

- Block of memory of size `eltsize` co-located with address `ptr`
- Accessed using traditional C semantics
- Leaves you at the new location

```
void * = mw_localmalloc(size_t S, void *p);
```

```
long * A = (long *) mw_localmalloc(sizeof(long), &B[2]);
```



Example: SAXPY with local allocation

```
#include <stdlib.h>
#include <cilk/cilk.h>
#include <memoryweb.h>
void saxpy4(long n, float a, float *x, float *y)
{
    for (long i = 0; i < n; i++)
        y[i] += a * x[i];
}
int main(int argc, char **argv)
{
    long num = atol(argv[1]); // number blocks
    long size = atol(argv[2]); // block size
    float aval = atof(argv[3]); // constant
    float **x = mw_malloc2d(num, size * sizeof(*x));
    float *y = mw_localmalloc(num * size * sizeof(*y), x[0]);
    for (long j = 0; j < num; j++)
        for (long i = 0; i < size; i++) {
            x[j][i] = j * size + i; y[j * size + i] = 0;
        }
    for (long i = 0; i < num; i++) {
        cilk_spawn_at (x[i]) saxpy4(size, aval, x[i], &y[i * size]);
    }
    cilk_sync;
}
```

- 2D Block allocation for input
- Output uses local allocation on particular node

Place output
on Node 0



Sample Program Execution: saxpy_loc_at.c

```
>>>>>>> /usr/local/emu/bin/emu-cc saxpy_loc_at.c -o saxpy_loc_at.mwx
>>>>>>> /usr/local/emu/bin/emusim.x --total_nodes 4 - saxpy_loc_at.mwx 4 8 5.0

SystemC 2.3.3-Accellera --- Mar 24 2021 16:05:40
Copyright (c) 1996-2018 by all Contributors,
ALL RIGHTS RESERVED
Start untimed simulation with local date and time= Thu Mar 25 22:40:38 2021

End untimed simulation with local date and time= Thu Mar 25 22:40:38 2021

>>>>>>> more saxpy_loc_at.vsf
Node ID: Outbound Migrations, Threads Created, Threads Died, Spawn Fails
0: 108, 1, 5, 0
1: 37, 1, 0, 0
2: 37, 1, 0, 0
3: 37, 1, 0, 0

NodeID.MspID: num_reads, num_writes, num_rmws
0.0: 131, 168, 5
1.0: 126, 206, 10
2.0: 97, 72, 154
3.0: 103, 156, 53
```

Spawns in
correct locations

More migrations due
to local allocation



Distributed Free

- Free data allocated by `mw_malloc2d` or `mw_malloc1dlong`

```
void mw_free(void *p);
```

- Free data allocated by `mw_localmalloc`

```
void mw_localfree(void *p);
```



Unit Summary: Data Distribution

- Cyclic data distribution
- Thread placement
- Block data distribution
- Varying parallelism just by changing data allocation, not algorithm code

Exercises:

Try different grain and block sizes,
check migrations and thread counts
Try different pointers to direct spawn
location
Try block allocation of structures
Use printf and see how migration
counts change

Extra Credit Answers:

Pointer index is a
multiple of number of
nodes, always on node 0





Data Distribution: Part II

Replicated allocation using memoryweb library

Global Replicated Data

- Instructs compiler to place an instance on each node
- Uses a “View 0” address that always gives local instance
- Must be a global variable
- Example Uses:
 - Constants
 - Copy on each node
 - All initialized to the same unchanging value
 - EX: PI, pointer to shared data structure
 - Local data
 - Copy on each node
 - May have different values
 - Use only when it does not matter which instance you access!
 - EX: random number table, pointer to local work queue

replicated long c = 3927883;



Dynamic Replicated Pointers

- Allocates a block on each node, returns replicated pointer
- Similar to using the replicated keyword
- Used when the size of the data structure is not known at compile time

```
long * mw_mallocrepl(size_t N);
```



Accessing Replicated Data Structures

- Return a pointer to the instance of a replicated data structure co-located with the destination pointer

```
void * mw_get_localto(void *src, void *dst);
```

- Return a pointer to the nth instance of a replicated data structure

```
void * mw_get_nth(void *src, unsigned n);
```




Example: Replicated Constant AX+Y

```
#include <stdlib.h>
#include <cilk/cilk.h>
#include <memoryweb.h>
long aval;
replicated long a;
void saxpy(long n, long a, long *x, long *y) {
    for (long i = 0; i < n; i++) y[i] += a * x[i];
}
int main(int argc, char **argv)
{
    long nth = atol(argv[1]); // number threads
    long size = atol(argv[2]); // array size
    aval = atol(argv[3]); // constant
    mw_replicated_init(&a, aval);
    long *x = mw_malloc1dlong(size);
    long *y = mw_malloc1dlong(size));
    for (long i = 0; i < size; i++) {
        x[i] = i; y[i] = 0;
    }
    long grain = size / nth; // elts per thread
    for (long i = 0, j = 0; i < nth; i++, j += grain)
        cilk_spawn saxpy(grain, &x[j], &y[j]);
    cilk_sync;
}
```

➤ All threads
access global
variable
instance on
current node

Constant is in
global replicated
variable



Sample Program Execution: saxpy_1d_repl.c

```
>>>>>>> /usr/local/emu/bin/emu-cc saxpy_1d_repl.c -o saxpy_1d_repl.mwx
>>>>>>> /usr/local/emu/bin/emusim.x --total_nodes 4 -- saxpy_1d_repl.mwx 4 32 5

SystemC 2.3.3-Accellera --- Mar 24 2021 16:05:40
Copyright (c) 1996-2018 by all Contributors,
ALL RIGHTS RESERVED
Start untimed simulation with local date and time= Thu Mar 25 22:40:38 2021

End untimed simulation with local date and time= Thu Mar 25 22:40:38 2021

>>>>>>> more saxpy_1d_repl.vsf
Node ID: Outbound Migrations, Threads Created, Threads Died, Spawn Fails
0: 8, 4, 5, 0
1: 8, 0, 0, 0
2: 8, 0, 0, 0
3: 8, 0, 0, 0

NodeID.MspID: num_reads, num_writes, num_rmws
0.0: 28, 25, 11
1.0: 75, 68, 27
2.0: 43, 26, 74
3.0: 9, 31, 1
```

➤ Same migration and thread pattern as spawn loop

➤ Memory accesses are slightly different

Extra Credit Question:

What would happen if global non-replicated variable were used instead?
(Answer at end of section)



Initializing Replicated Data with Value

- Initializes each instance of replicated data structure to the same value

```
void = mw_replicated_init(long *p, long v);
```

```
replicated long N;  
int main()  
{  
    ...  
    long n_elements = compute_n_elements();  
    mw_replicated_init(&N, n_elements);  
    ...  
}
```



Initializing Replicated Data using Node ID

- Initializes each instance of the replicated data structure using the result of the user-defined function `init_func(n)` where `n` is the node number

```
void mw_replicated_init_multiple  
(long *p, long (*f)(long));
```

```
replicated long B;  
long init_func(long nid) {  
    return nid * 5;  
}  
int main()  
{  
    ...  
    mw_replicated_init_multiple(&B, init_func);  
    ...  
}
```



Initializing Replicated Data using Pointer

- Initializes each instance of replicated data structure using the user-defined function $f(&obj, n)$ where $&obj$ is the address of the replicated data structure and n is the node number

```
void mw_replicated_init_generic  
    (long *p, void (*f)(void *, long));  
  
replicated struct info { long x; long y; } info;  
void init_info(void *obj; long nid) {  
    struct info *i = (struct info*) obj;  
    i->x = node;  
    i->y = 5*node + 4;  
}  
int main()  
{  
    ...  
    mw_replicated_init_multiple(&info, init_info);  
    ...  
}
```



Unit Summary: Data Replication

- Replicated global variables
- Allocation, initialization and use within threads
- Affects on migrations and memory accesses

Exercises:

Try other replicated data structures

Try more complex initializations

Measure effects of replicated vs. non-replicated data

Extra Credit Answer:

Non-replicated variable
will cause more
migrations in code

