

Outline

This presentation covers the following topics

➤ Lucata Workflow

- *X86 Debugging*
- *Simulation*
- *Hardware*

*Slides originally developed
by Janice McMahon,
Lucata Corporation*



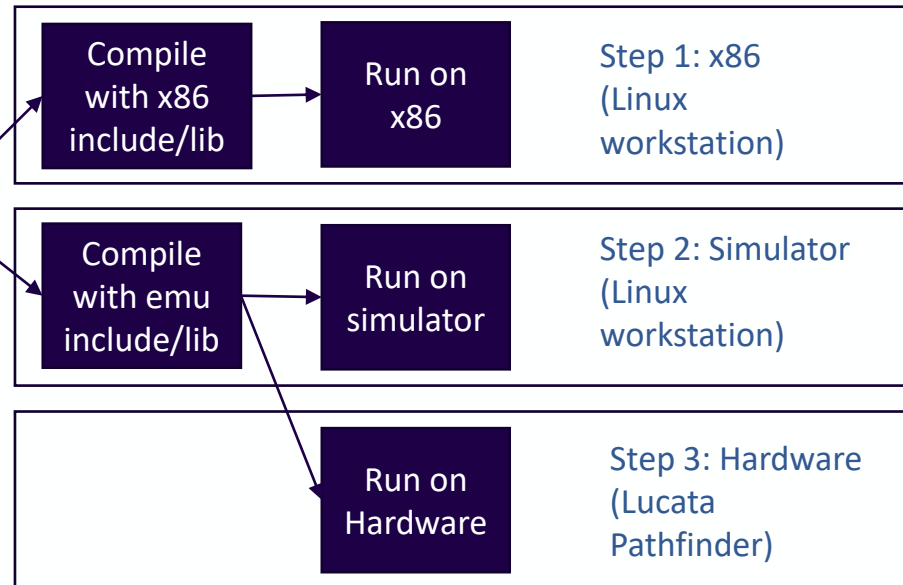


Debugging

Workflow step 1: x86 execution for verification of correctness

Single program

```
#ifdef X86
#include "memoryweb_x86.h"
#else
#include "memoryweb.h"
#endif
// rest of C/C++ Cilk program
```

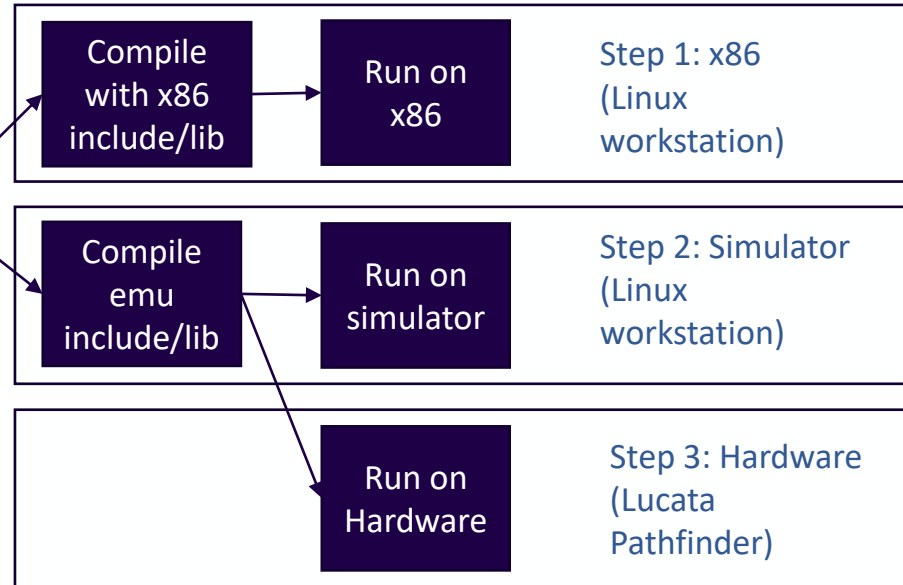


Software Development Workflow

Single program

```
#ifdef X86
#include "memoryweb_x86.h"
#else
#include "memoryweb.h"
#endif
// rest of C/C++ Cilk program
```

- Only difference is include file
- Program uses intrinsics, mw_malloc functions for distributed data
- X86 version mimics single node with single core



Development steps should be done in sequence
Lucata hardware platform used only to run final code
All tools run on Linux workstation

Workflow Step 1

Single program

```
#ifdef X86
#include "memoryweb_x86.h"
#else
#include "memoryweb.h"
#endif
// rest of C/C++ Cilk program
```

Compile
with x86
include/lib

"exe"
file

Run on
x86

Step 1: x86

- Uses standard Linux compiler requiring Cilk support (GCC v5-v7 or OpenCilk)

Example Linux commands:

```
> clang -DX86 -I /usr/local/emu/x86/include main.c -o  
main -L /usr/local/emu/x86/lib -lemu_c_utils  
> main
```

→ Produces executable file

→ Runs program on x86

- Compile with special paths (flags to compiler)
- Could use standard Linux tools (editor, debugger, profiler)
- Program is built and run the same manner as any C program

**Verify correct operation of
parallel Lucata program**

Debug in x86 mode

- Provides cross-compilation for Emu codes on x86
- Use for rapid building and testing of codes before deployment to Emu architecture
- Treats system as single node with multiple Cilk threads
- Requires Cilk support in x86 compiler
- Use x86 library and include paths when building



Sample Program Execution: intrs_hook.c

```
#include <cilk/cilk.h>
#ifdef X86
#include <stdio.h>
#include <memoryweb_x86.h>
#include <emu_c_utils.h>
#else
#include <memoryweb.h>
#include <emu_c_utils/emu_c_utils.h>
#endif
#include <emu_c_utils/emu_c_utils.h>
...
```

- Uses different include files
- Uses different directories for includes and libraries

We recommend using OpenCilk

Program is executed as in standard Linux manner

```
$ opencilk-2.0.1/bin/clang -I /usr/local/emu/x86/include/emu_c_utils
-DX86 intrs.c -L /usr/local/emu/x86/lib -l emu_c_utils -o intrs

$ ./intrs
cycles = 13056

$ opencilk-2.0.1/bin/clang -I
/usr/local/emu/x86/include/emu_c_utils -DX86 intrs_hook.c -L
/usr/local/emu/x86/lib -l emu_c_utils -o intrs_hook

$ ./intrs_hook
{"region_name":"example","time_ms":0.69,"ticks":347008}
time (ms) = 0.694016
```



Unit Summary: Debugging

- Code modifications for x86 execution
- Building and running in x86 mode

Exercises:

- Re-build all examples for x86
- Use standard Linux tools (gprof, gdb, etc.) on x86 executable





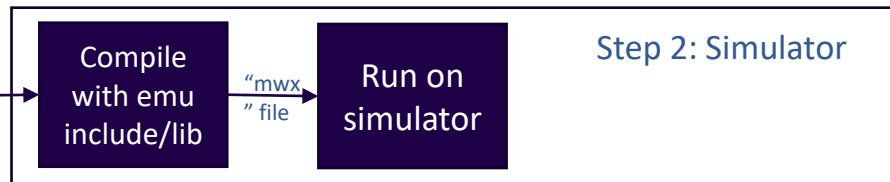
Simulation

Workflow step 2: implementation study using simulator

Workflow Step 2

Single program

```
#ifdef X86
#include "memoryweb_x86.h"
#else
#include "memoryweb.h"
#endif
// rest of C/C++ Cilk program
```



- Use Lucata compiler, linker and simulator on Linux platform
- Simulator produces thread migration and memory usage statistics
- Used with profiler for visualization of those statistics

Example Linux commands:

```
> emu-cc -o main.mwx main.c -lemu_c_utils
> emusim.x -- main.mwx
> emusim_profile dir - main.mwx
```

Runs simulator and produces image files from statistics that can be viewed in a browser

Produces "mwX" executable file
Simulator mimics execution of mwX and produces output files with detailed statistics

Verify correctness of Emu Cilk program
Architecture modeling and study
Understand parallel performance characteristics for *small data sets*

Simulation Modes

- Untimed mode produces summary statistics
 - Runs a functional rather than a timed simulation
- Timed mode produces detailed statistics
 - Wealth of detail included in output files (*.vsf, *.cdc, *.msp)
 - Entered via function call in code (used in hooks functions)
 - Specifies that ALL code AFTER the call will be included in detailed timing measurements
 - Required for CLOCK intrinsic in the simulator; produces NOOP on hardware
 - Used to generate raw data for profiler images



Simulator Options for Program Control

➤ Simulator Control

- Initialization of memory
- Run untimed (functional)
- Maximum simulation time
- Sampling interval
- Return values
- Output monitor period
- Help
- Output file base

--initialize_memory
--profile functional
--max_sim_time
--timing_sample_interval
--*return_value*
--output_monitor_period
-h, --help
-o, --base_ofile

➤ Machine Configuration

- Memory per Node
- Total Nodes

-m, --log2_memory_per_node
-n, --total_nodes



Simulator Options for Profiling Execution

- Execution information printed to screen
 - List all spawn, quit, migrate operations to screen
--{untimed}_short_trace
 - Thread ID, Node ID, TPC, Memory and Register effects for every instruction executed
--verbose_isa
 - Verbose thread information
--verbose_tid
- Instruction counts per function (.uis file)
--output_instruction_count
- Queue statistics (.tqd file)
--capture_timing_queues

Used by
profiler to
generate data
for images



Example: Short Trace

```
>>>>>>> /usr/local/emu/bin/emusim.x --total_nodes 4 --untimed_short_trace -- saxpy_ld.mwx 4 32 5
Start untimed simulation with local date and time= Thu Apr  1 17:16:11 2021
```

```
TID0 NODE 0 SPAWN CHILD TID1 FUNC @main @cycle 2634
TID1 NODE 0 MIGRATE DEST 1 FUNC @main.outline_.otd1 @cycle 2673 TPC 0x800024ef Inst LDE
TID0 NODE 0 SPAWN CHILD TID2 FUNC @main @cycle 2687
TID1 NODE 1 MIGRATE DEST 2 FUNC @main.outline_.otd1 @cycle 2693 TPC 0x800024ef Inst LDE
TID1 NODE 2 MIGRATE DEST 3 FUNC @main.outline_.otd1 @cycle 2713 TPC 0x800024ef Inst LDE
TID2 NODE 0 MIGRATE DEST 1 FUNC @main.outline_.otd1 @cycle 2726 TPC 0x800024ef Inst LDE
TID1 NODE 3 MIGRATE DEST 0 FUNC @main.outline_.otd1 @cycle 2733 TPC 0x800024ef Inst LDE
TID0 NODE 0 SPAWN CHILD TID3 FUNC @main @cycle 2740
TID2 NODE 1 MIGRATE DEST 2 FUNC @main.outline_.otd1 @cycle 2746 TPC 0x800024ef Inst LDE
TID1 NODE 0 MIGRATE DEST 1 FUNC @main.outline_.otd1 @cycle 2753 TPC 0x800024ef Inst LDE
TID2 NODE 2 MIGRATE DEST 3 FUNC @main.outline_.otd1 @cycle 2766 TPC 0x800024ef Inst LDE
TID1 NODE 1 MIGRATE DEST 2 FUNC @main.outline_.otd1 @cycle 2773 TPC 0x800024ef Inst LDE
TID3 NODE 0 MIGRATE DEST 1 FUNC @main.outline_.otd1 @cycle 2779 TPC 0x800024ef Inst LDE
TID2 NODE 3 MIGRATE DEST 0 FUNC @main.outline_.otd1 @cycle 2786 TPC 0x800024ef Inst LDE
TID0 NODE 0 SPAWN CHILD TID4 FUNC @main @cycle 2793
TID1 NODE 2 MIGRATE DEST 3 FUNC @main.outline_.otd1 @cycle 2793 TPC 0x800024ef Inst LDE
TID3 NODE 1 MIGRATE DEST 2 FUNC @main.outline_.otd1 @cycle 2799 TPC 0x800024ef Inst LDE
TID2 NODE 0 MIGRATE DEST 1 FUNC @main.outline_.otd1 @cycle 2806 TPC 0x800024ef Inst LDE
TID1 NODE 3 MIGRATE DEST 0 FUNC @main @cycle 2813 TPC 0x8000238d Inst ADDM
TID1 NODE 0 DIED NODE 0 FUNC @main @cycle 2817
TID3 NODE 2 MIGRATE DEST 3 FUNC @main.outline_.otd1 @cycle 2819 TPC 0x800024ef Inst LDE
TID2 NODE 1 MIGRATE DEST 2 FUNC @main.outline_.otd1 @cycle 2826 TPC 0x800024ef Inst LDE
TID4 NODE 0 MIGRATE DEST 1 FUNC @main.outline_.otd1 @cycle 2832 TPC 0x800024ef Inst LDE
TID3 NODE 3 MIGRATE DEST 0 FUNC @main.outline_.otd1 @cycle 2839 TPC 0x800024ef Inst LDE
TID2 NODE 2 MIGRATE DEST 3 FUNC @main.outline_.otd1 @cycle 2846 TPC 0x800024ef Inst LDE
TID4 NODE 1 MIGRATE DEST 2 FUNC @main.outline_.otd1 @cycle 2852 TPC 0x800024ef Inst LDE
TID3 NODE 0 MIGRATE DEST 1 FUNC @main.outline_.otd1 @cycle 2859 TPC 0x800024ef Inst LDE
TID2 NODE 3 MIGRATE DEST 0 FUNC @main @cycle 2866 TPC 0x8000238d Inst ADDM
TID2 NODE 0 DIED NODE 0 FUNC @main @cycle 2870
TID4 NODE 2 MIGRATE DEST 3 FUNC @main.outline_.otd1 @cycle 2872 TPC 0x800024ef Inst LDE
TID3 NODE 1 MIGRATE DEST 2 FUNC @main.outline_.otd1 @cycle 2879 TPC 0x800024ef Inst LDE
TID4 NODE 3 MIGRATE DEST 0 FUNC @main.outline_.otd1 @cycle 2892 TPC 0x800024ef Inst LDE
TID3 NODE 2 MIGRATE DEST 3 FUNC @main.outline_.otd1 @cycle 2899 TPC 0x800024ef Inst LDE
TID4 NODE 0 MIGRATE DEST 1 FUNC @main.outline_.otd1 @cycle 2912 TPC 0x800024ef Inst LDE
TID0 NODE 0 DIED NODE 0 FUNC @main @cycle 2918
TID3 NODE 3 MIGRATE DEST 0 FUNC @main @cycle 2919 TPC 0x8000238d Inst ADDM
TID3 NODE 0 DIED NODE 0 FUNC @main @cycle 2923
TID4 NODE 1 MIGRATE DEST 2 FUNC @main.outline_.otd1 @cycle 2932 TPC 0x800024ef Inst LDE
TID4 NODE 2 MIGRATE DEST 3 FUNC @main.outline_.otd1 @cycle 2952 TPC 0x800024ef Inst LDE
TID4 NODE 3 MIGRATE DEST 0 FUNC @main @cycle 2972 TPC 0x8000238d Inst ADDM
TID4 NODE 0 DIED NODE 0 FUNC @Exit @cycle 3623
End untimed simulation with local date and time= Thu Apr  1 17:16:11 2021
```

Trace thread
movement
throughout
program
execution to
verify expected
migration
patterns



Object Dump for Debugging

```
>>>>>>> /usr/local/emu/bin/emu-cc saxpy_1d.c -o saxpy_1d.mwx  
/usr/local/emu/bin/gossamer64-objdump -D saxpy_1d.mwx
```

```
saxpy_1d.mwx:      file format elf64-gossamer64
```

```
Disassembly of section .text:
```

```
0000000040001000 <@saxpy>:  
  40001000: 80002000: ETD      2  
  40001001: 80002002: BCTGT    0x8000200a  
  40001003: 80002006: JMP      0x8000203f  
  40001005: 8000200a: LSR0  
  40001007: 8000200d: DTE      6  
  40001008: 8000200f: ETD      6  
  40001009: 80002011: SLLC     3  
  4000100b: 80002015: DPETA    4  
  4000100c: 80002018: LDE      7  
  4000100e: 8000201b: ETD      6  
  4000100f: 8000201d: SLLC     3  
  40001011: 80002021: DPETA    5  
  40001012: 80002024: ETD      3  
  40001013: 80002026: MULTE    7  
  40001015: 80002029: ADDM  
  40001016: 8000202b: ETA      6  
  40001017: 8000202d: AAIMB    1  
  40001018: 80002030: ATE      6  
  40001019: 80002032: ETD      2  
  4000101a: 80002034: XORE     6  
  4000101c: 80002037: BCTDZ    0x8000203f  
  4000101e: 8000203b: JMP      0x8000200f  
  40001020: 8000203f: JMPE     1  
  
0000000040001021 <@main>:  
  40001021: 80002042: ETD      0  
  40001022: 80002044: DTD2  
  40001023: 80002046: LSR3  
  40001025: 80002049: DTE      0  
  40001026: 8000204b: LIT16    128=0x0080  
  40001029: 80002051: SILL     0x0  
  4000102c: 80002057: SILL     0x0
```

Uses Emu
version of
standard Linux
objdump utility

Thread program counter
can be correlated with
simulator output or
simulator exception
message to pinpoint
errors

Simulator Error Exception

```
>>>>>>> /usr/local/emu/bin/emusim.x --total_nodes 4 -- saxpy_1d.mwx

SystemC 2.3.3-Accellera --- Mar 24 2021 16:05:40
Copyright (c) 1996-2018 by all Contributors,
ALL RIGHTS RESERVED
Start untimed simulation with local date and time= Thu Apr 1 17:38:35 2021
```

```
Translating address == 0
UNTIMED SIMULATION on NODE[0] TID[0] generated exception
ThreadID=0
HW ThreadID=0x1
Thread using HW ThreadID
ThreadletState=Active
ThreadletException=1=Address
ExecutionType=7
Current Instruction:
```

```
80004eb2 LD8A: iToken=162 iLength=3 nibbles=b3e000
```

```
Threadlet TCB Data:
TCB.(TPC)=(0x80004eb2) (32 bits each)
TCB.(D,D2)=(1,1) (one bit each)
TCB.(A,A2)=(1,1) (one bit each)
TCB.(TS,TSDATA)=(0,0x0) (two bits, four bits)
TCB.AID=0x1 (8 bits)
TCB.MaxDepth=0xff (8 bits)
TCB.Priority=0x1 (8 bits)
TCB.(NaN,U,V,CB,N,Z)=(0, 0, 0, 0, 0, 0)
TCB.M=0 (one bit)
TCB.DB=1 (one bit)
```

```
Threadlet Data Registers
```

```
A: 0x0=0
```

```
A2: 0x180000080000a90=108086393204378256
```

```
Format: signed decimal, unsigned decimal, hex
D: 2147503792, 2147503792, 0x80004eb0
D2: 0, 0, 0x0
```

**Manufactured error
(arguments missing)**

```
E[0] (Live): 108086393204378256, 108086393204378256, 0x180000080000a90
E[1] (Live): 2147492262, 2147492262, 0x800021a6
E[2] (Live): 0, 0, 0x0
E[3] (Live): 108086393204375568, 108086393204375568, 0x1800000800000010
E[4] (Live): 1152921504606848016, 1152921504606848016, 0x10000000000000410
E[5] (Live): 36028797018978840, 36028797018978840, 0x8000000000003a18
E[6] (Live): 36028797018977280, 36028797018977280, 0x8000000000003400
E[7] (Live): 36028797018977288, 36028797018977288, 0x80000000000003408
E[8] (Live): 1, 1, 0x1
E[9] (Live): 1040, 1040, 0x410
E[10] (Live): 108086395351859176, 108086395351859176, 0x1800000fffffe8
E[11] (Live): 108086393204378248, 108086393204378248, 0x180000080000a88
E[12] (Live): 0, 0, 0x0
E[13] (Live): 108086393204375568, 108086393204375568, 0x1800000800000010
E[14] (Live): 0, 0, 0x0
E[15] (Live): 0, 0, 0x0
```

```
Other Useful Data
```

```
Fence Counter=0
```

```
Source Node=0
```

```
Dest Node=0
```

```
End untimed simulation with local date and time= Thu Apr 1 17:38:35 2021
```

```
halt_data_dump function has been removed
```

```
.halt_data_dump function has been removed
```

```
.halt_data_dump function has been removed
```

```
.halt_data_dump function has been removed
```

```
.GENERATED EXCEPTION
```

**Find TPC in object
dump to locate error**

Unit Summary: Simulation

- Many options for controlling simulator to exercise a variety of machine configurations
- Many options for producing detailed information about program execution to aide in understanding performance
- Used in conjunction with Linux utilities to pinpoint errors at low level of code

Exercises:

- Run previous examples with timed/untimed mode, examine data file contents, study migration patterns via short-trace.





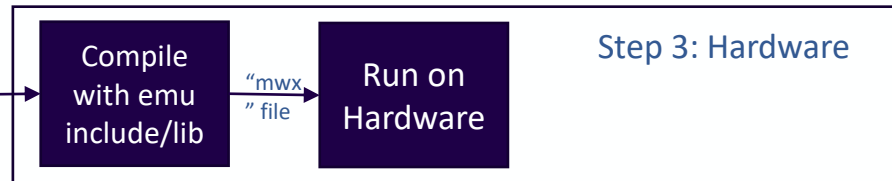
Hardware

Workflow step 3: execution for timing measurement

Workflow Step 3

Single program

```
#ifdef X86
#include "memoryweb_x86.h"
#else
#include "memoryweb.h"
#endif
// rest of C/C++ Cilk program
```



- Use Lucata compiler and linker on Linux platform
- Executable must be copied to Lucata machine over LAN and run on that machine (i.e., cross-compiled)
- Execution time can be measured but no statistics gathered

Example Linux commands:

```
> emu-cc -o main.mwx main.c -lemu_c_utils
> scp main.mwx LUCATA:
> ssh LUCATA
LUCATA> emu_multimode_exec main.mwx
```

Produces "mw" executable file

Copy mw to Lucata machine over network

Run command executed on Lucata machine

Run and measure program on Lucata machine

Pathfinder Configuration

➤Single-node Execution

- Program runs on a single node
- Uses all Gossamer cores on that node
- Users can work independently on different nodes

➤Multi-node Execution

- Program runs on full system
- Can access all nodes (all Gossamer cores)
- Single user



Pathfinder Hardware Execution

- Compile programs on Host then scp to Pathfinder
- Single-node Execution
 - Launched on node using `emu_handler_and_loader`
- Multi-node Execution
 - Launched on node 0 using `emu_multinode_exec`



Program Execution Utilities

- Load program and data to all nodes
- Launch initial thread into the system
- Monitor the system exception queue and handle system services until a thread quits or an exception occurs
- Terminate by issuing a checkpoint to clear the system and dump any remaining threads
- Print information to log files for each thread that quits, exits, generates an exception, or is checkpointed
- Return the program's return value



Slurm Extensions for the Pathfinder

- Georgia Tech has developed a Slurm workflow for the Pathfinder to simplify this process
- No need to scp data/programs!
 - All data is shared via network-mounted storage
- Use sbatch to launch jobs for single-node, single-chassis, or multi-chassis jobs

```
#Show the real-time status of the Pathfinder nodes.  
%sinfo --federation -M pathfinder
```

				CLUSTER:	pathfinder
PARTITION	AVAIL	TIMELIMIT	NODES	STATE	NODELIST
rg-pathfinder*	up	180-00:00:	1	alloc	c2n1
			31	idle	c0n[0-7],c1n[0-7],c2n[0,2-7],c3n[0-7]

```
#Run the Saxpy command with Slurmemu_handler_and_loader saxpy-1d-workflow.mwx 8 128 5.0  
sbatch -M pathfinder -q single-node --wrap "emu_handler_and_loader 0 0 -- saxpy-1d-workflow.mwx 8 128 5.0"
```

