



# Hardware Performance Counters

Janice McMahon  
September 2022

# Motivation for Hardware Performance Counters

---

- The Lucata workflow allows for many opportunities to check functional correctness via x86 execution and simulation
- However, understanding system performance may require more detailed insight into the hardware performance counters that can be used.
- There are 5 types of possible counters
  - Node Counters
  - MSP Counters
  - GC Cluster Counters
  - GC Counters
  - SRIO Counters



# Function for Hardware Performance Counters

---

➤ Single library function controls counters

**void `lu_profile_perfcntr`(perfcntr\_e op, char \*str);**

➤ “Op” argument determines operation to be performed

- PFC\_CLEAR: initialize/clear counters
- PFC\_START: start the counters without reading
- PFC\_READ: output current counter values
- PFC\_STOP: stop and read counters

➤ “Str” argument determines string for output file

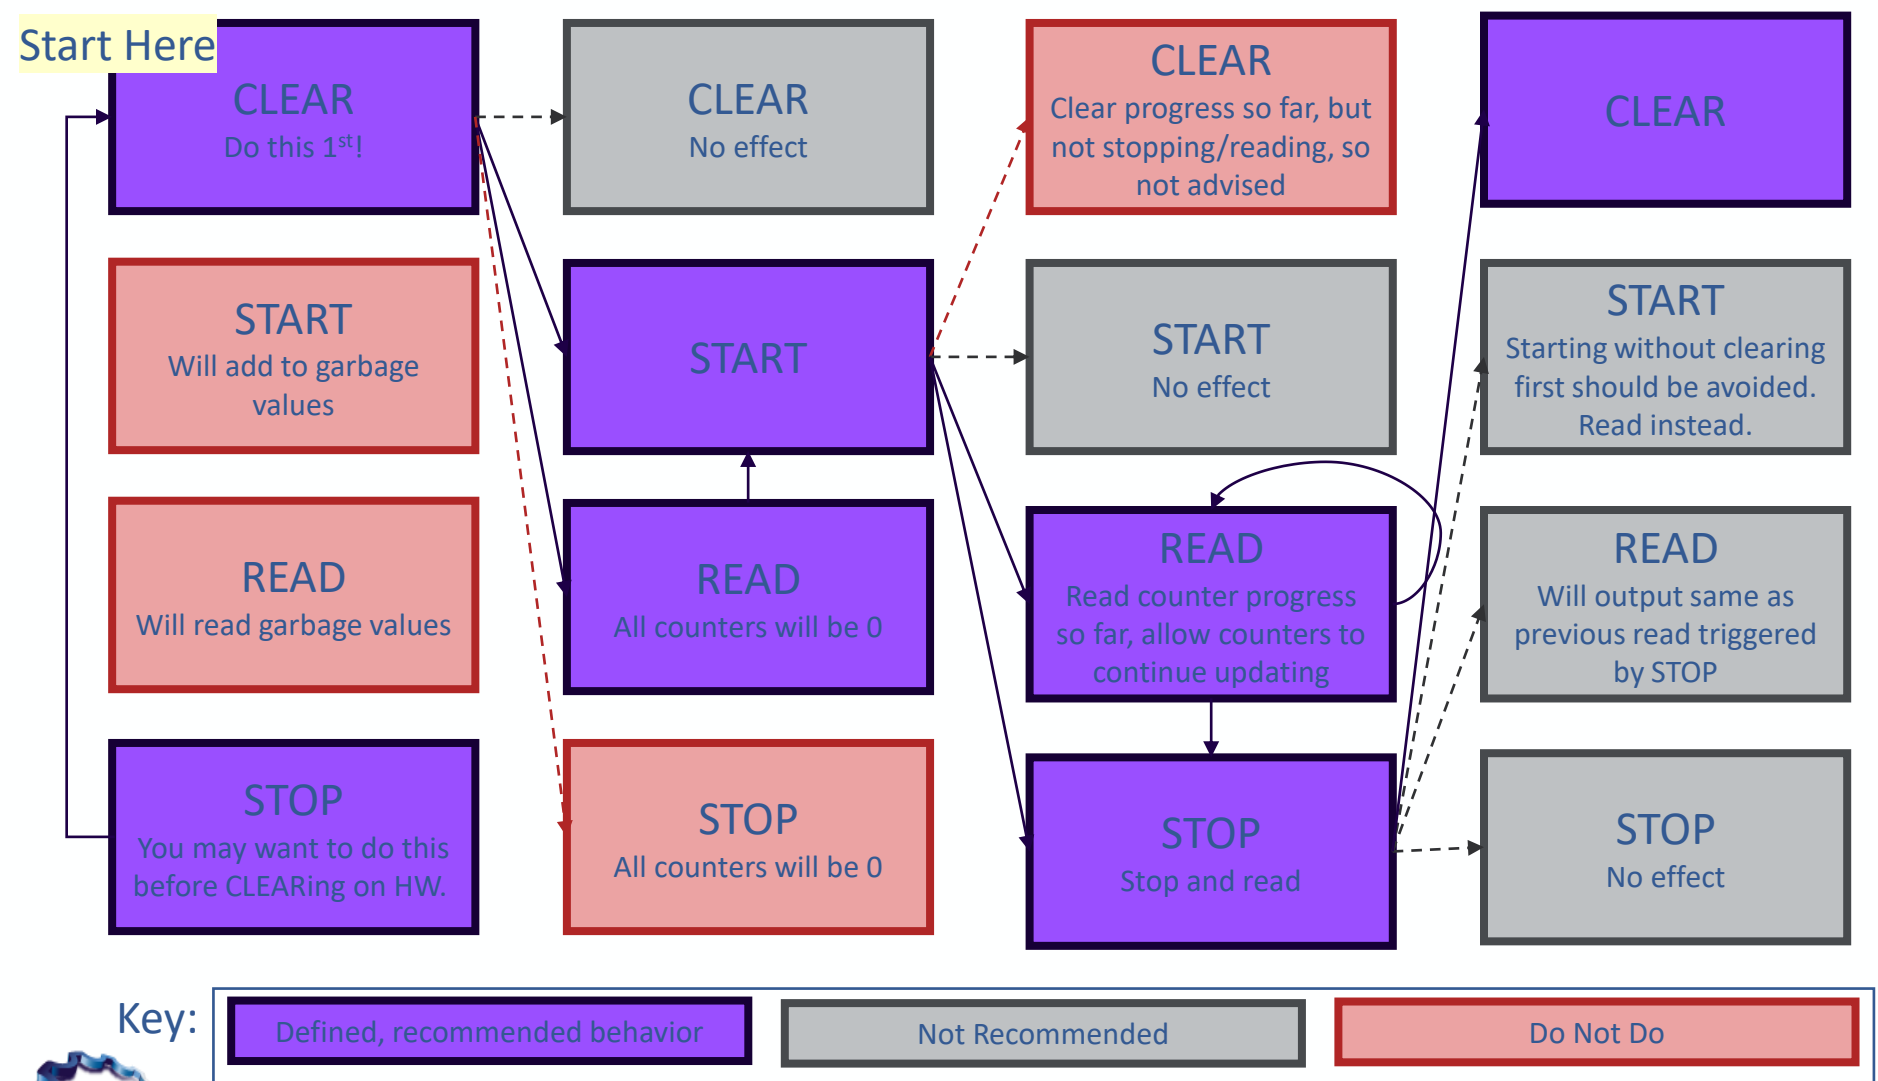
- Used as key for JSON object
- Should be unique for all calls to profiler

➤ Use only during serial portion of code

➤ Must be used according to rules for well-defined behavior



# Performance Counter Usage Flowchart



# Performance Counters: example

```
starttiming();
lu_profile_perfcntr(PFC_CLEAR, "SIMPLE ARRAY SUM CLEARING COUNTERS");
lu_profile_perfcntr(PFC_READ, "SIMPLE ARRAY SUM READING COUNTERS AFTER CLEARING");
lu_profile_perfcntr(PFC_START, "SIMPLE ARRAY SUM STARTING COUNTERS");5

// Spawn child thread to sum values in A
sumA = cilk_spawn sum(A, epn);

// Call function to sum values of B (in current/parent thread)
sumB = sum(B, epn);

// Wait for spawned thread to complete
cilk_sync;

printf("A = %ld, B = %ld\n", sumA, sumB);
lu_profile_perfcntr(PFC_STOP, "SIMPLE ARRAY SUM STOPPING COUNTERS AT END");
return 0;
```

Initialize and start counters in  
beginning, serial portion of program

Stop and read counters  
at end of program



# Performance Counter Output

```
1 {
2   "SIMPLE ARRAY SUM CLEARING COUNTERS": "N/A",
3   "SIMPLE ARRAY SUM READING COUNTERS AFTER CLEARING": {
4     "System_Parameters": {
5       "Total_Nodes":2,
6       "MSPs_Per_Node":8,
7       "GC_Clusters_Per_Node":3,
8       "GCs_Per_GC_Cluster":8
9     },
10    "MSP_Counters": {
11      "MSP_memory_reads": {
12        "Node0": {
13          "MSP0":0,
14          "MSP1":0,
15          "MSP2":0,
16          "MSP3":0,
17          "MSP4":0,
18          "MSP5":0,
19          "MSP6":0,
20          "MSP7":0
21        },
22        "Node1": {
23          "MSP0":0,
24          "MSP1":0,
25          "MSP2":0,
26          "MSP3":0,
27          "MSP4":0,
28          "MSP5":0,
29          "MSP6":0,
30          "MSP7":0
31        }
32      },
33    }
```

```
1 "SIMPLE ARRAY SUM STARTING COUNTERS": "N/A",
2 "SIMPLE ARRAY SUM STOPPING COUNTERS AT END": {
3   "System_Parameters": {
4     "Total_Nodes":2,
5     "MSPs_Per_Node":8,
6     "GC_Clusters_Per_Node":3,
7     "GCs_Per_GC_Cluster":8
8   },
9   "MSP_Counters": {
10     "MSP_memory_reads": {
11       "Node0": {
12         "MSP0":49,
13         "MSP1":23,
14         "MSP2":144,
15         "MSP3":153,
16         "MSP4":7,
17         "MSP5":70,
18         "MSP6":61,
19         "MSP7":109
20       },
21       "Node1": {
22         "MSP0":35,
23         "MSP1":5,
24         "MSP2":2,
25         "MSP3":0,
26         "MSP4":0,
27         "MSP5":0,
28         "MSP6":0,
29         "MSP7":0
30       }
31     ...
```

Output  
in JSON  
format

.hpc file  
created  
automatically  
by simulator

# Hardware Execution

---

- Node managing software creates log files for each node in JSON format
- Counter output .hpc file is created by Python script that aggregates hardware log files

```
1 usage: emu_pc_concat.py [-h] -f PC_FILE [-d OUTPUT_DIR]
2
3 optional arguments:
4 -h, --help show this help message and exit
5 -f PC_FILE, --pc_file PC_FILE
6             mn_exec_pc.PID.log file name
7 -d OUTPUT_DIR, --output_dir OUTPUT_DIR
8             path to directory to save .hpc file if you don't want
9             to save it to default location (path to
10             mn_exec_pc.PID.log file)
```



# Performance Counter Graphs

- Graph images are created by Python script that operates on .hpc output file

```
1 usage: make_hpc_plots.py [-h] -f HPC_FILE [-l LOG_FILE] [-n] [-s]
2
3 optional arguments:
4   -h, --help show this help message and exit
5   -f HPC_FILE, --hpc_file HPC_FILE
6               .hpc file name
7   -l LOG_FILE, --log_file LOG_FILE
8               .out file name
9   -n, --no_graphs only make csvs, no graphs
10  -s, --stop_reads stop parsing JSON after a specific read
```





# Graph Directory Structure

```
— simpleArraySum_hwperfctrs_30-11-2021_15:37:39
  — SIMPLE_ARRAY_SUM_READING_COUNTERS_AFTER_CLEARING
    — GC_Cluster_Counters
    — GC_Counters
    — MSP_Counters
    — SRIO_In_Counters
    — SRIO_Out_Counters
  — SIMPLE_ARRAY_SUM_STOPPING_COUNTERS_AT_END
    — GC_Cluster_Counters
    — GC_Counters
    — MSP_Counters
    — SRIO_In_Counters
    — SRIO_Out_Counters
```

# Performance Counter Comparison

- Counter reads from different versions of program compared by Python script that operates on .hpc output files

```
1 usage: compare_hwperfcntr_jsons.py [-h] -f FILES [FILES ...] -d SAVE_DIR [-c]
2
3 optional arguments:
4   -h, --help show this help message and exit
5   -f FILES [FILES ...], --files FILES [FILES ...]
6       list of .hpc files to compare
7   -d SAVE_DIR, --save_dir SAVE_DIR
8       where to save the comparison graphs
9   -c, --compare_all Use this flag to compare all sets instead of choosing
10      which sets of counters
```

## Usage:

```
./compare_hwperfcntr_jsons.py -f simpleArraySum_hwperfcntrs_diff_ele.hpc
simpleArraySum_hwperfcntrs.hpc -d example_dir
```



# Performance Counter Polling

---

- Frequent sampling of performance counters on hardware to gather data on system changes over time
- Activated via flag to hardware execution command
- Raw data .csv file and graphs created by Python script executing on hardware log file with polling data



# Graph Generation from Polling Data

```
1 usage: emu_analyze_pr.py [-h] -f PR_FILE [-d OUTPUT_DIR] [-p PREFIX] [-n]
2                        [-c NUM_INTERVALS] [-v] [-l] [-s] [-b] [-i INPUT_CSV]
3
4 optional arguments:
5   -h, --help show this help message and exit
6   -f PR_FILE, --pr_file PR_FILE
7                        mn_exec_pr.PID.log file name
8   -d OUTPUT_DIR, --output_dir OUTPUT_DIR
9                        path to directory to save output files to if you don't
10                       want to save it to default location (path to
11                       mn_exec_pr.PID.log file)
12   -p PREFIX, --prefix PREFIX
13                        Add prefix to graph directory name and graph names
14   -n, --no_graphs only make csvs, no graphs
15   -c NUM_INTERVALS, --choose_intervals NUM_INTERVALS
16                        choose which intervals to graph
17   -v, --no_csv do not create csv
18   -l, --no_max_line do not add max threads line to threads graphs
19   -s, --stacked_gc_plot
20                        graph GC CONTEXTS as stacked area plot
21   -b, --gc_subplots graph GC CONTEXTS as one subplot per GC
22   -i INPUT_CSV, --input_csv INPUT_CSV
23                        path to existing csv, don't have to re-create. CAN
24                        SAVE A LOT OF TIME
```



# Example: Array Sum

```
...
int main()
{
    allocate_arrays(); // Allocate arrays using malloc2d
    initialize_arrays(); // Initialize arrays

    starttiming();
    lu_profile_perfcntr(PFC_CLEAR, "CLEAR COUNTERS");
    lu_profile_perfcntr(PFC_READ, "READ COUNTERS AFTER CLEAR");
    lu_profile_perfcntr(PFC_START, "START COUNTERS");

    sumA = cilk_spawn sum(A, epn); // Spawn child thread for A
    sumB = sum(B, epn); // Sum values of B in parent thread
    cilk_sync; // Wait for spawned thread to complete
    printf("A = %ld, B = %ld\n", sumA, sumB);
    lu_profile_perfcntr(PFC_STOP, "STOP COUNTERS AT END");
}
```

Clear,  
read,  
then  
start  
counters

Read  
after  
clear;  
values  
are zero

Stop counters at end;  
implicit read of values



# Sample Program Execution: Array\_Sum.c

```
>>>>>>> emu-cc ArraySum_hpc.c -o ArraySum_hpc.mwx -l memoryweb
>>>>>>> emusim.x ArraySum_hpc.mwx

SystemC 2.3.3-Accellera --- Sep  7 2022 09:15:59
Copyright (c) 1996-2018 by all Contributors,
ALL RIGHTS RESERVED
Start untimed simulation with local date and time= Wed Sep 14 12:49:27 2022

[WARN]: In /home/tdysart/Toolchains/22.R2-Sept7/llvm-
cilk/mwsim/src/timsim.cpp:sc_main:590 Calling READCNTRS when not profiling!
Timed simulation starting...
A = 32, B = 64
End untimed simulation with local date and time= Wed Sep 14 12:49:31 2022

Info: /OSCI/SystemC: Simulation stopped by user.
>>>>>>> ls ArraySum_hpc.*
ArraySum_hpc.c  ArraySum_hpc.hpc  ArraySum_hpc.mwx
ArraySum_hpc.cdc  ArraySum_hpc.mps
>>>>>>> make_hpc_plots.py -f ArraySum_hpc.hpc
Find all graphs in: ./ArraySum_hpc_14-09-2022_12:49:55
The last hpc call to analyze will be 1
Program called lu_profile_perfcntnr with message: SIMPLE ARRAY SUM READING
COUNTERS AFTER CLEARING
Generating Graphs for [SIMPLE ARRAY SUM READING COUNTERS AFTER CLEARING]...
/home/tdysart/Toolchains/22.R2-Sept7/llvm-cilk/lucata-
22.R2.Sept7/bin/make_hpc_plots.py:177: UserWarning: Tight layout not
applied. The bottom and top margins cannot be made large enough to
accommodate all axes decorations.
    plt.tight_layout()
Program called lu_profile_perfcntnr with message: SIMPLE ARRAY SUM STOPPING
COUNTERS AT END
Generating Graphs for [SIMPLE ARRAY SUM STOPPING COUNTERS AT END]...
Stopping here after read 1
hpc_file_name_base: ArraySum_hpc.hpc
>>>>>>> ls ArraySum_hpc_14-09-2022_12\49\55/*
ArraySum_hpc_14-09-2022_12:49:55/ArraySum_hpc.hpc.csv

'ArraySum_hpc_14-09-
2022_12:49:55/SIMPLE_ARRAY_SUM_READING_COUNTERS_AFTER_CLEARING':
GC_Cluster_Counters  MSP_Counters
```

**Counter  
output file  
created**

**Graphs created  
in subdirectory**

- Simulator creates .hpc file automatically
- Profile images generated after simulation by script
- Subdirectory created for each read of counters in program


```
'ArraySum_hpc_14-09-
2022_12:49:55/SIMPLE_ARRAY_SUM_STOPPING_COUNTERS_AT_END':
GC_Cluster_Counters  GC_Counters  MSP_Counters
```




# Hybrid BFS: MSP Counters I

2 Nodes, Graph Scale 5


## DRAM Writes




## Memory Writes




## Block Writes




## DRAM Reads



## Memory Reads




## Block Reads



## Hybrid BFS: MSP Counters II

# Atoms




## Contexts Dequeued




## Contexts Enqueued



## Response Words



## Sub-Word Writes



2 Nodes  
Graph Scale 5






# Hybrid BFS: GC Counters I


## Node 0

### Internal Spawns




## Node 1

### Internal Spawns




2 Nodes  
Graph Scale 5

### Spawn Failures




### Spawn Failures



# Hybrid BFS: GC Counters II


Node 0

Stall Cycles




Node 1

Stall Cycles




2 Nodes  
Graph Scale 5

Total Successful Spawns




Total Successful Spawns




# Hybrid BFS: GC Cluster Counters I

2 Nodes, Graph Scale 5


## Atomics




## Block Reads




## Block Writes




## Contexts Loaded



## Enq Packets




## Memory Reads




# GC Cluster Counters II


## Memory Writes



## Migrations Out



## Request Transactions



## Response Transactions



## Service Requests




2 Nodes  
Graph Scale 5




# BFS: SRIO In I


## Acks




## Block Read Requests



## Block Write Requests




## EnqTSR Requests




2 Nodes  
Graph Scale 5

# BFS: SRIO In II


## Read Requests



## Read Responses




## Remotes




2 Nodes  
Graph Scale 5

# BFS: SRIO Out I

Acks




Block Read Requests




2 Nodes  
Graph Scale 5

Block Write Requests




EnqTSR Requests




# BFS: SRIO Out II


## Read Requests



## Read Responses



## Remotes




2 Nodes  
Graph Scale 5



# BFS Comparison: Beamer vs. Migrating Threads


## MSP Memory Writes

RUN\_BFS\_TRIAL\_1\_STOP-G5\_N2\_migrating\_threads/hybrid\_bfs\_bfs\_no\_worklist\_hpc\_T256.hpc  
RUN\_BFS\_TRIAL\_1\_STOP-G5\_N2\_beamer\_hybrid/hybrid\_bfs\_bfs\_no\_worklist\_hpc\_T256.hpc

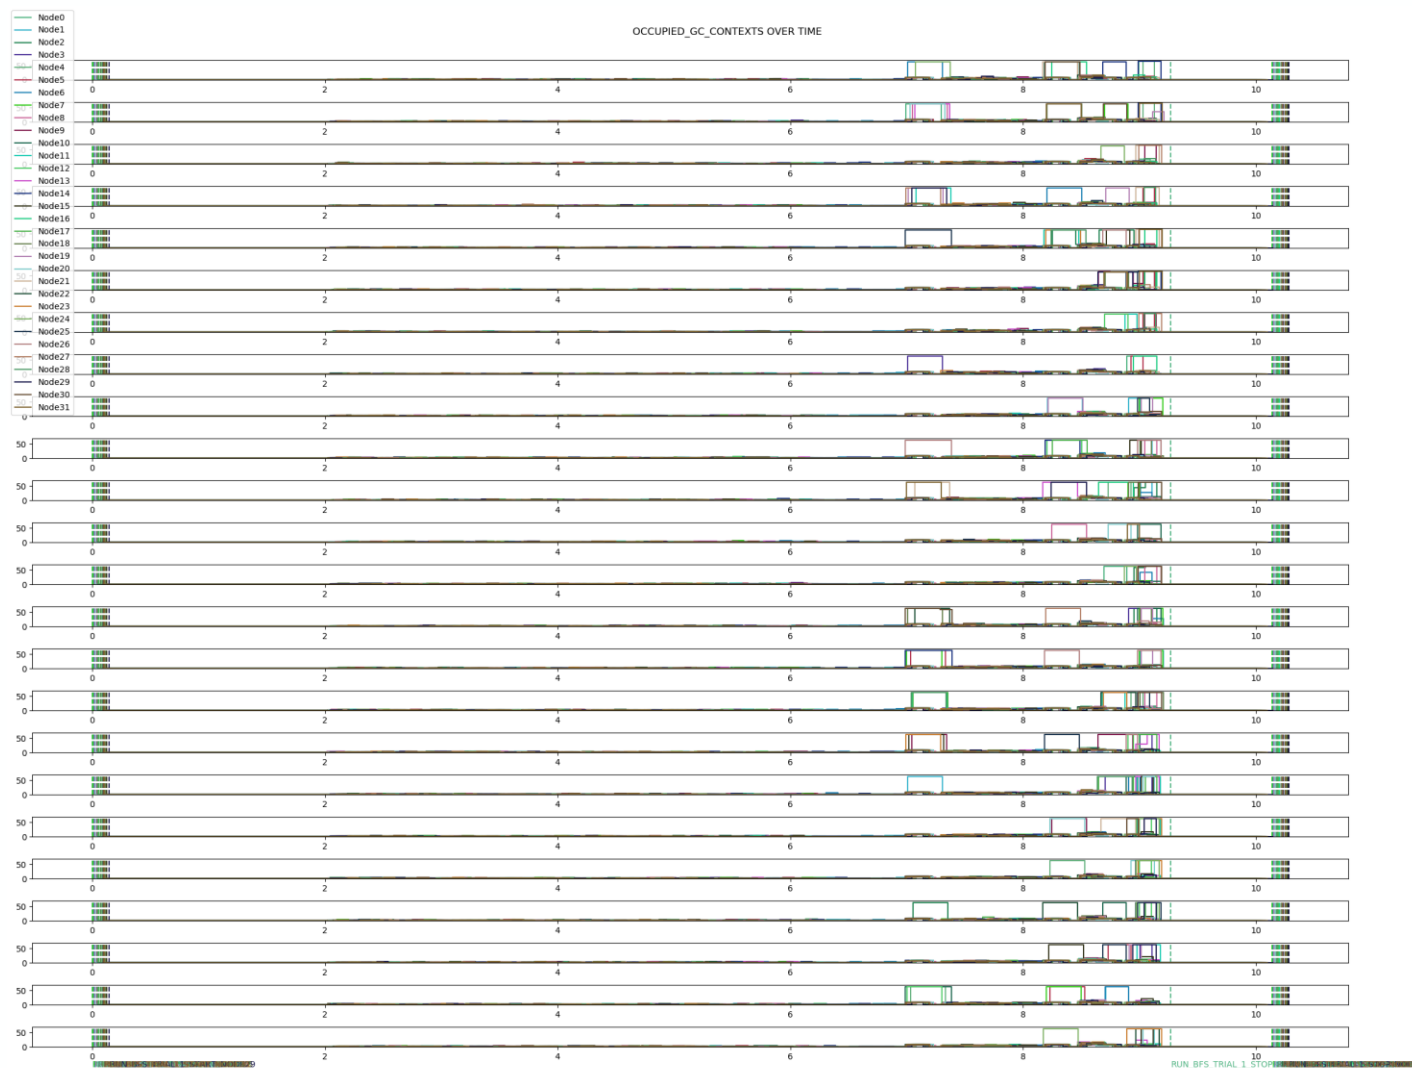


2 Nodes  
Graph Scale 5

# BFS Polling: Number of Threads over Time




# BFS Polling: GC Contexts over Time



32 Nodes  
100 ms  
Graph  
Scale 28

# BFS Polling: GC Contexts Stacked



32 Nodes  
100 ms  
Graph  
Scale 28

# Unit Summary: Performance Counters

---

- Library call inserted in code
- Graphs generated by script
- Example graphs shown

## Exercises:

Try reading counters with more complex codes (extra spawns to create failures, extra memory operations)

Try inserting multiple read commands in code

Try comparing versions and running on hardware with polling

