

Outline

This presentation covers the following topics

- Session 1 – Cilk and the Lucata API
 - Basic programming
 - Data distribution
- Session 2 – Lucata Workflow
 - X86 Debugging
 - Simulation
 - Hardware
- Session 3 – Measuring Performance
 - Timing Hooks
 - Profiling
- Session 4 – Coding Optimizations
 - Machine-specific coding
 - Parallel computation
- **Section 5**
 - ***Advanced topics***

*Slides originally developed
by Janice McMahon, Lucata
Corporation*





Advanced Topics

Patterns for spawning threads

Balancing Parallelism and Overhead

- Number of threads vs work per thread
 - Enough parallelism to keep the cores busy and mask migrations
 - Enough work per thread to offset thread overhead
- Target ~64 threads per core
 - Larger systems and/or more migrations may require more threads to offset those in transit
 - Maximum 512 threads per node



Common Issues


- Hotspots/Extra Migrations
 - Shared data found on single node
 - Spawns not distributed
- Wave parallelism
 - All threads move in lock step, not distributed
- Unbalanced/limited network usage
 - Sequential : only uses channel to neighbor e.g. BFS
 - Other non-uniform communication e.g. GUPS



"Wave" parallelism

```
#define N 8
#pragma grainsize = 4
cilk_for (long i=0; i<N; i++) {
    E[i]->a = i;
    E[i]->b = 0;
}
```

- Accessing a striped array in order will cause a migration on every visit
- Starting all threads on node 0 utilizes only 1 node at a time



Spawning Static Thread Teams

- Spawn a “team” of threads at each node
- Each thread has a “home” node
 - Threads may stay entirely local
 - Threads may migrate away then return for more work
- If work varies greatly, a work queue often performs better than assigning N elements to each thread
 - Automatically load balances
 - Reduces cost of spawn and sync by creating a single set of threads
 - Atomic increment used to grab next unit of work



Dynamic Spawning

- Traverse a data structure and spawn based on characteristics of the structure
- Example: BFS
 - Spawn a team of threads to process vertices
 - Dynamically spawn additional threads to process edges in parallel based on size of the edge list
 - Number of edges at each vertex is unknown and may vary greatly



Efficient Spawning

- Distribute spawns across the nodes then spawn additional local threads
- Recursively spawn threads in a tree-like fashion (for large # threads)
 - Allows parallel spawning rather than sequential
 - Reduces hotspot at a single node if spawning across multiple nodes
 - Used in `cilk_for`



Cilk_for Spawn Tree Parallelization

```
cilk_for (i=0; i < E; i++) {c[i] = a[i]+ b[i];};
```

spawn two threads

```
(i=0; i < E/2; i++);
```

```
(i=E/2; i < E; i++);
```

Each spawns 2 more threads

```
(i=0; i < E/4; i++);
```


```
(i=E/4; i < E/2; i++);
```

```
(i=3*E/4; i < E; i++);
```

```
(i=E/2; i < 3*E/4; i++);
```

...

Recursive vs. Linear Spawn




Linear migration pattern
with sequential spawns


```
void linear() {  
    for (long i = 0; i < NUM_NODES(); i++)  
    { // rvar is a replicated variable  
        long *p = mw_get_nth(&rvar, i);  
        cilk_spawn_at (p) node_function();  
    }  
}
```

Tree migration pattern
with parallel spawns

```
void recurse(long s, long e) {  
    while (1) {  
        long count = e - s;  
        if (count <= 1) break;  
        long m = s + count / 2;  
        long *p = mw_get_nth(&rvar, m);  
        cilk_spawn_at (p) recurse(m, e);  
        e = m;  
    } // assert (s == NODE_ID());  
    node_function(s);  
}
```

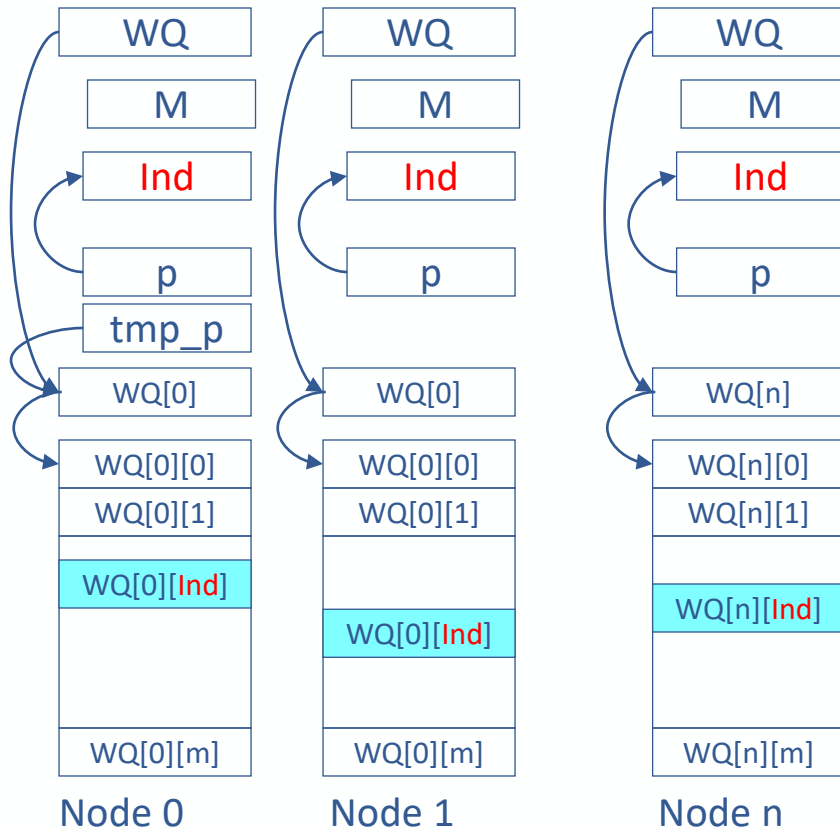


Work Queue in each Node



- Multiple threads are spawned on each node
- Each thread grabs the next work item from the queue for processing
- Next work item is designated by an index in node memory
- Threads grab items via an atomic increment of the index in that node
 - Uses the `ATOMIC_ADDMS` intrinsic
 - Increments memory and returns previous value

Data Structures for Work Queue



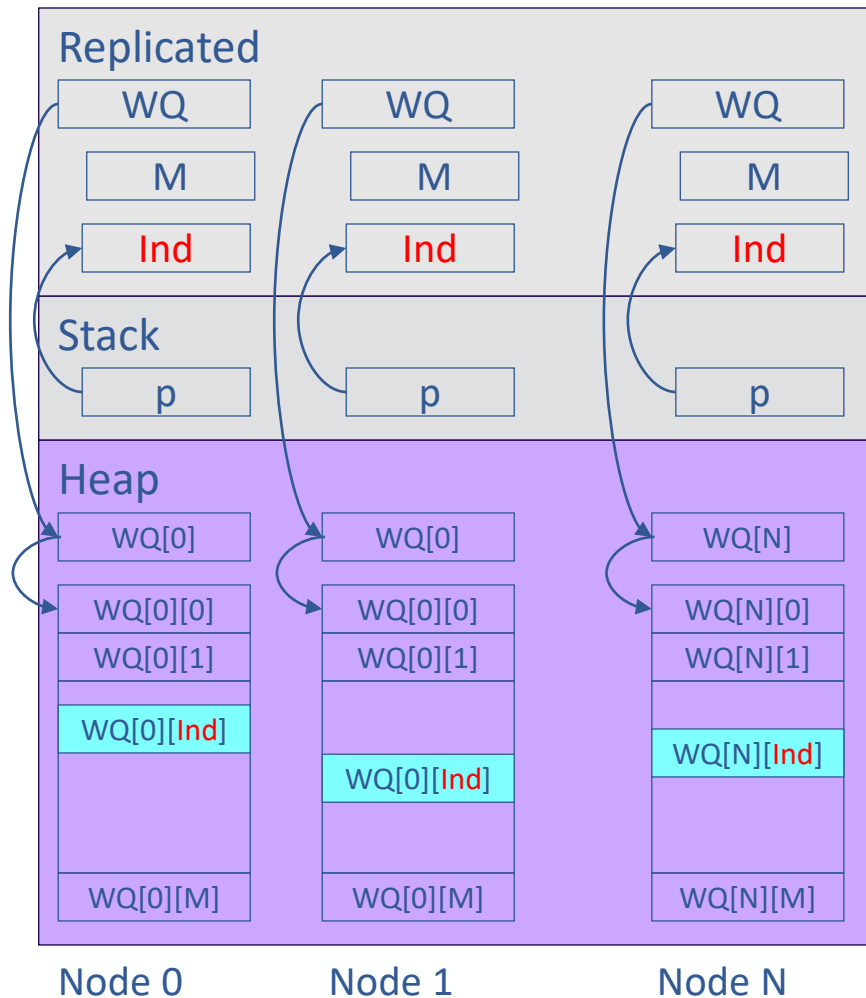
Pointer is constant and can be replicated to save migrations

```

replicated dtype **WQ;
replicated long M;
replicated long Ind;
main() {
    long N = NUM_NODES(); // 10 items/node
    mw_replicated_init(&M, 10);
    dtype **tmp_p = // work item type
        mw_malloc2d(N, M * sizeof(dtype));
    mw_replicated_init(&WQ, tmp_p);
    recurse(0, N);
}

void node_function(long nnum) {
    // assert (nnum == NODE_ID());
    long *p = mw_get_nth(&Ind, nnum);
    *p = 16; // 16 threads per node
    for (long i = 0; i < 16; i++)
        cilk_spawn worker(p, i, nnum);
}
    
```

Worker Thread Function



Each thread grabs the next work item using **ATOMIC_ADDMS**; done when all items processed

```
void worker(long *p, long off, long n) {
    // assert n == NODE_ID();
    while (off < M) {
        ...
        // process item WQ[n][off]
        ...
        off = ATOMIC_ADDMS(*p, 1);
    }
}
```



Replicated variable complexity

- Using replicated pointers reduces number of migrations over both argument values and global variable with function calls
- With thread spawns, arguments are packaged into thread state, no migration for access
- Register spills can cause extra migrations, eliminating the performance gain from using replicated variables
- Use of `noinline` keyword essential for preventing unwanted migrations due to register spills

