

# Outline

---

This presentation covers the following topics

- Session 1 – Cilk and the Lucata API
  - Basic programming
  - Data distribution
- Session 2 – Lucata Workflow
  - X86 Debugging
  - Simulation
  - Hardware
- Session 3 – Measuring Performance
  - Timing Hooks
  - Profiling
- **Session 4 – Coding Optimizations**
  - ***Machine-specific coding***
  - ***Parallel computation***
- Section 5
  - Advanced topics

*Slides originally developed  
by Janice McMahon, Lucata  
Corporation*





## Machine-specific Coding

*Intrinsics including atomics, remotes, compare-and-swap*

# Atomic Arithmetic Operations

- Atomic arithmetic operations at the memory
- Four distinct “flavors” of each atomic that specify the value returned and resulting memory value

**long `ATOMIC_<INST>`(long \*p, long D);**  
**<INST>: ADD/AND/OR/XOR/MAX/MIN**

Mnemonic	Return Value	New Memory Value
<op>	Result	Unchanged
<op>S	Result	Orig. D Value
<op>M	Result	Result
<op>MS	Orig. Mem. Value	Result



# ATOMIC\_ADD Examples

## ➤ **ATOMIC\_ADD(A, D)**

- Reads value at A
- Adds D
- Returns result
- Mem[A] is unchanged

## ➤ **ATOMIC\_ADDS(A, D)**

- Reads value at A
- Adds D
- Returns results
- Mem[A] = D

## ➤ **ATOMIC\_ADDM(A, D)**

- Reads value at A
- Adds D
- Returns result
- Mem[A] = result

## ➤ **ATOMIC\_ADDMS(A, D)**

- Reads value at A
- Adds D
- Returns original Mem[A]
- Mem[A] = result



# Remote Arithmetic Operations

- Significantly reduces migrations by performing atomic updates to memory without migrating the thread
- Sends only the data and operation to be performed
  - Consumes less than half the bandwidth of a typical thread migration
- Does not return a value
- Returns an ACK, may be turned off
- Remotes issued by a thread to the same location guaranteed to complete in order

```
void REMOTE_<INST>(long *p, long D);  
<INST>: ADD/AND/OR/XOR/MAX/MIN
```



# Memory Fence on Remote Update

---

- Used to wait for all remote updates to be acknowledged
- Prevents thread from continuing until all outstanding acknowledgements have been received
- Implicit FENCE before migration, thread suspend, etc.

**void FENCE();**



# Swap Operations

## ➤Atomic swap operation

- Replace the contents of memory at A with D
- Return the original contents of memory at A

**long `ATOMIC_SWAP`(long \*A, long D);**

## ➤Atomic compare and swap operation

- Compare the contents of memory at A to "cmp"
- If they match, swap "new" into A
- Always return original contents of memory at A

**long `ATOMIC_CAS`(long \*A, long new, long cmp);**



## Other Operations

---

### ➤Population count

- Counts number of 1s in the word referenced by "ptr"
- Adds this value to sum and returns the result

**long POPCNT(long sum, long \*ptr);**

### ➤Priority

- Computes the 6-bit priority encode on "value"
- i.e., bit position of highest numbered non-zero bit

**unsigned long PRIORITY(unsigned long value);**





# Thread Management

---

- Resize the thread to carry only the live registers
  - Used by compiler or programmer before a possible migration to reduce thread size

**void RESIZE();**

- Reschedule the thread to the end of the Run Queue
  - Allow a new thread to be scheduled in the core
  - Can be used to minimize the impact of busy wait

**void RESCHEDULE();**



## System Query

- Return system time on local node (clock ticks)

**unsigned long CLOCK();**

- Return the current thread ID

**unsigned long THREAD\_ID();**

- Return the current Node ID

**unsigned long NODE\_ID();**

- Return the number of nodes in the current system

**unsigned long NUM\_NODES();**

- Return the number of bytes per node

**unsigned long BYTES\_PER\_NODE();**

- Return the number of Gossamer cores per node


**unsigned long GCS\_PER\_NODE();**



# Example: Distributed Array of Integers

```
#define N 256 // elements per node
long ** A = mw_malloc2d(NUM_NODES(), N * sizeof(long));
for (long j = 0; j < NUM_NODES(); j++)
    for (long i = 0; i < N; i++)
        A[j][i] = i;
```


Uses intrinsic for number  
of nodes in system



# Allocate Replicated Data Structure

```
struct stats {  
    long count;  
    long max;  
    long min;  
};  
replicated struct stats s;
```


Copy of “stats”  
structure on  
each node at  
the same offset



# Initialize Replicated Data Structure

```
for (long i = 0; i < NUM_NODES(); i++) {  
    struct stats * si = mw_get_nth(&s, i);  
    si->count = 0;  
    si->max = 0;  
    si->min = LONG_MAX;  
}
```

Initialize using  
mw\_get\_nth




# Update Replicated Data Structure

```
cilk_for (long i = 0; i < N; i++) {  
    long score = scoreval(A[i]);  
    if (score != 0) {  
        ATOMIC_ADDM(&(s.count), 1);  
        ATOMIC_MAXM(&(s.max), score);  
        ATOMIC_MINM(&(s.min), score);  
    }  
}
```

Loops over array;  
updates instance  
on node co-located  
with thread


Uses intrinsics for atomic  
operations on shared data



# Reduce Replicated Data Structure

```
unsigned long count = 0;
unsigned long max = 0;
unsigned long min = LONG_MAX;
for (long i = 0; i < NUM_NODES(); i++) {
    struct stats * si = mw_get_nth(&s, i);
    count += si->count;
    if (si->max > max) max = si->max;
    if (si->min < min) min = si->min;
}
```


Reduce  
over  
replicated  
instances  
on each  
node



# Example: Intrinsic

```
int main(int argc, char **argv)
{
    long **A = distr_array(); // create 2d array
    initialize(); // initialize replicated data
    unsigned long start = CLOCK(); // clock start
    // update replicated data
    for (long I = 0; I < NUM_NODES(); i++) {
        cilk_spawn_at (A[i]) update(A[i]);
    }
    cilk_sync;
    // reduce replicated data
    long count = 0;
    long max = 0;
    long min = LONG_MAX;
    for (long i = 0; i < NUM_NODES(); i++) {
        struct stats *si = mw_get_nth(&s, i);
        count += si->count;
        if (si->max > max) max = si->max;
        if (si->min < min) min = si->min;
    }
    long total = CLOCK() - start; // total cycles
    printf("cycles %lu count %ld max %ld min %ld\n",
        total, count, max, min);
}
```

Use intrinsic  
to measure  
clock cycles





# Sample Program Execution: intrs.c

```
>>>>>>> /usr/local/emu/bin/emu-cc intrs.c -o intrs.mwx
>>>>>>> /usr/local/emu/bin/emusim.x --total_nodes 4 -- intrs.mwx

SystemC 2.3.3-Accellera --- Mar 24 2021 16:05:40
Copyright (c) 1996-2018 by all Contributors,
ALL RIGHTS RESERVED
Start untimed simulation with local date and time= Sun Mar 28 23:33:05 2021

cycles = 4599
End untimed simulation with local date and time= Sun Mar 28 23:33:05 2021

>>>>>>> more intrs.vsf
Node ID: Outbound Migrations, Threads Created, Threads Died, Spawn Fails
0: 6, 64, 68, 0
1: 3, 64, 63, 0
2: 3, 64, 63, 0
3: 3, 64, 63, 0

NodeID.MspID: num_reads, num_writes, num_rmws
0.0: 411, 594, 184
1.0: 958, 558, 1362
2.0: 402, 600, 146
3.0: 447, 529, 296
```

- Thread creations distributed over nodes
- Migrations only for initial thread spawn
- All accesses are local to replicated variables



# Unit Summary: Machine-Specific Coding

- Use of intrinsics for allocating data and distributing data precisely
- Use of intrinsics to measure clock cycles
- Use of atomics to update node-local statistics shared among multiple threads
- Use of intrinsics to perform reductions over replicated variables

## Exercises:

Measure other codes using intrinsics

Use atomics to implement more complex structures such as distributed queues

Use remote intrinsics to avoid migrations altogether





## Parallel Computation

*Local, distributed, and chunked arrays in C/C++ utilities*

# Parallel Computation with C Utilities

---

- Functions for local, distributed striped, and distributed chunked
- Programmer supplies function pointer, library spawns threads across system for parallel slides
- Grain size (work per threadlet) can be computed automatically via system settings



# Functions for Local Arrays

---

- Apply function to range in parallel

**void emu\_local\_for(...);**

- Initialize array in parallel

**void emu\_local\_for\_set\_long(...);**

- Copy array in parallel

**void emu\_local\_for\_copy\_long(...);**

- Copy a range of bytes

**void emu\_memcpy(...);**

- Choose grain size based on array size

**static inline long LOCAL\_GRAIN\_MIN(...);**



# Local Arrays: example

```
void worker(long begin, long end, va_list args)
{
    long *x = va_arg(args, long*);
    long b = va_arg(args, long);
    for (long i = begin; i < end; ++i) x[i] += b;
}
```

Worker function called on each array slice over iteration space defined by <begin, end>

Spawn parallel threads within a node:

```
long *x = malloc(1024 * sizeof(long));
long begin = 0;
for (long i = 0; i < 1024; i += grain)
{
    long end = MIN(i + grain, 1024);
    cilk_spawn worker(i, end, x, 5);
}
```

Grain size is number of elements to assign to each thread; computed from total size and number of threads

Requires number of threads and grain size to be program variables within application code



# Local Arrays: C utilities

```
void worker(long begin, long end, va_list args)
{
    long *x = va_arg(args, long*);
    long b = va_arg(args, long);
    for (long i = begin; i < end; ++i) x[i] += b;
}
```

Same worker function

Spawn parallel threads within a node:

```
long *x = malloc(1024 * sizeof(long));
emu_local_for(0, 1024, LOCAL_GRAIN_MIN(1024, 64),
              worker, x, 5);
```

Grain size parameter  
is minimum number  
of elements to assign  
to each thread

Built-in function computes  
grain size suitable for  
processing all items on a single  
node given a lower bound



# Functions for Distributed (striped) Arrays

---

- Distributed parallel for loop

```
void emu_1d_array_apply(...);
```

- Distributed parallel reduction

```
void emu_1d_array_reduce_sum(...);
```

- Choose grain size based on array size

```
static inline long GLOBAL_GRAIN_MIN(...);
```





# Striped Arrays: example

```
void worker(long *array, long begin, long end, va_list args)
{
    long *x = array;
    long b = va_arg(args, long);
    for (long i = begin; i < end; i += NUM_NODES()) x[i] += b;
}
```

Worker function called on each array slice over iteration space defined by <begin, end> with stride NUM\_NODES()

## Spawn parallel threads within a node:

```
long * x = malloc1dlong(1024 * sizeof(long));
long begin = 0;
for (long j = 0; j < NUM_NODES(); j++) {
    for (long i = 0; i < 1024; i += grain) {
        long end = MIN(i + grain, 1024);
        cilk_spawn worker(&x[j], i, end, 5);
    }
}
```

Grain size is number of elements to assign to each thread; computed from total size and number of threads

Requires number of nodes, number of threads and grain size to be program variables within application code; requires distributed spawn tree to be efficient



# Striped Arrays: C utilities

```
void worker(long *array, long begin, long end, va_list args)
{
    long *x = array;
    long b = va_arg(args, long);
    for (long i = begin; i < end; i += NUM_NODES()) x[i] += b;
}
```

Same worker function

Spawn parallel threads within a node:

Grain size parameter is minimum number of elements to assign to each thread

```
long *x = malloc1dlong(1024);
emu_1d_array_apply(x, 1024,
    GLOBAL_GRAIN_MIN(1024, 64), worker, 5);
```

Built-in function computes grain size suitable for processing all items on the entire system; handles distributed spawn inside function hidden from programmer



# Functions for Distributed (chunked) Arrays

- Allocate and return replicated pointer

**emu\_chunked\_array \* emu\_chunked\_array\_replicated\_new(...);**

- Free replicated pointer

**void emu\_chunked\_array\_replicated\_free(...);**

- Initialize chunked array

**void emu\_chunked\_array\_replicated\_init(...);**

- Deallocate chunked array

**void emu\_chunked\_array\_replicated\_deinit(...);**

- Pointer to i-th element

**static inline void \* emu\_chunked\_array\_index(...);**

- Number of elements

**long emu\_chunked\_array\_size(...);**



## Functions for Distributed (chunked) Arrays (cont.)

---

- Distributed parallel for loop

**void emu\_chunked\_array\_apply(...);**

- Initialize array in parallel

**void emu\_chunked\_array\_set\_long(...);**

- Distributed parallel reduction

**long emu\_chunked\_array\_reduce\_sum(...);**

- Scatter from local array to distributed array

**void emu\_chunked\_array\_from\_local(...);**

- Gather from distributed array to local

**void emu\_chunked\_array\_to\_local(...);**



# Chunked Arrays: example

```
void worker(long *x, long begin, long end, va_list args)
{
    long b = va_arg(args, long);
    for (long i = begin; i < end - begin; ++i) x[i] += b;
}
```

## Spawn parallel threads within a node:

```
long **x = mw_malloc2d(1024, sizeof(long));
long begin = 0;
for (long j = 0; j < NUM_NODES(); j++) {
    for (long i = 0; i < 1024; i += grain) {
        long end = MIN(i + grain, 1024);
        cilk_spawn worker(&x[j][0], i, end, 5);
    }
}
```

Worker function called on each array slice of chunked array over iteration space defined by <begin, end>

Grain size is number of elements to assign to each thread; computed from total size and number of threads

Requires number of nodes, number of threads and grain size to be program variables within application code; requires replicated variables and distributed spawn to be efficient



# Chunked Arrays: C utilities

```
void worker(emu_chunked_array *array, long begin, long end, va_list args)
{
    long *x = emu_chunked_array_index(array, begin);
    long b = va_arg(args, long);
    for (long i = begin; i < end - begin; ++i) x[i] += b;
}
```

Worker function uses  
chunked array

Spawn parallel threads within a node:

```
emu_chunked_array *x =
    emu_chunked_array_replicated_new(1024, sizeof(long));
emu_chunked_array_apply(x, GLOBAL_GRAIN_MIN(1024, 64),
    worker, 5);
```

Grain size parameter  
is minimum number  
of elements to assign  
to each thread

Built-in function computes grain size suitable  
for processing all items on the entire system;  
hides distributed spawn from programmer



# Unit Summary: Parallel Computation

---

- C Utilities for parallel programming
- Local, distributed, and chunked arrays

## Exercises:

Re-implement saxpy using distributed, chunked arrays

